

PETadex: A comprehensive database of plastic-degrading enzyme sequence diversity

Dennis Zhu¹, Thomas Quigley¹, Artem Babaian¹

¹Donnelly Centre for Cellular and Biomolecular Research, University of Toronto, Toronto, ON, Canada.

Contributing authors: a.babaian@utoronto.ca;

Abstract

Introduction

Plastics are a cornerstone of modern life, improving the accessibility and affordability of necessities such as medical supplies and food packaging. They are indispensable in much of the Global South, where alternative materials are scarce and prohibitively expensive [cite]. The same traits responsible for these benefits, low production cost and durability, have also made plastics a persistent, large-scale pollutant: over 400 million tonnes are produced annually, and 13 billion tonnes have accumulated since the mid-20th century [plastic production citation here]. About half of this plastic waste has been recycled or incinerated (6.3 billion tonnes), while the remaining 6.7 billion tonnes are unaccounted for in the environment [plastic pollution citation here] (Fig. 1A). This unaccounted-for waste exceeds the mass of all animals on Earth (an estimated 2.5 billion tonnes [cite animal biomass]) by 2.6-fold (Fig. 1A).

Unaccounted-for plastic waste degrades into micro- and nanoplastics (MNPs), which permeate Earth’s ecosystems, enter food chains and bioaccumulate in humans [Thompson 2024 citation here]. MNPs should be considered toxic pollutants: polystyrene MNPs fed to mice cause focal disruption of the gut epithelial barrier and inflammatory cytokine release [cite].

Despite the relative novelty of plastics, some microorganisms have adapted to break them down using plastic-degrading enzymes (PDEs). In 2016, Yoshida *et al.* isolated *Ideonella sakaiensis*, a bacterium that uses polyethylene terephthalate (PET) as its major carbon and energy source, and identified the enzyme responsible, IsPETase [cite Yoshida 2016]. PDEs such as IsPETase represent a promising solution to plastic pollution: they break plastics down into their constituent monomers, which can then be repolymerized into new plastic products or repurposed into biomass (Fig. 1B).

Existing databases of PDEs are limited in scope and diversity. The plastics-active enzymes database (PAZy) contains 495 functionally verified plastic-active enzymes [cite PAZy], most of them homologs of leaf-branch compost cutinase (LCC) [cite] or IsPETase. Sequence-based mining is a well-established method for discovering new enzymes: Danso *et al.* (2018) [cite] identified 504 candidates by searching UniProtKB and NCBI-nr, and a further 349 by mining 133 metagenomes. More recently, Erickson *et al.* (2022) [cite] screened 87 hot-spring metagenomes (~38 million sequences), characterizing 74 putative enzymes, testing 51 for activity and identifying 23 novel plastic-active enzymes. These surveys have searched only a small fraction of the sequencing data now publicly available.




Fig. 1 Plastic waste and its enzymatic degradation. (A) Mass of plastics compared with Earth’s major groups of life (1 square = 1 billion tonnes). Of 13 Gt of plastics, 6.3 Gt has been recycled or incinerated and 6.7 Gt remains in landfill or the environment, 2.6-fold the mass of all animals (2.5 Gt). [update Fig. 1A graphic: label should read “recycled or incinerated”, not “recycled”] (B) Plastic-degrading enzymes as a route out of the plastics crisis: IsPETase hydrolyses PET to the monomer BHET, then to TPA and EG, which feed biomass or recycling.

In previous work, we built Logan [cite], an assembly of the entire Sequence Read Archive (SRA), 27.3 million metagenomic datasets, and searched it with 213 characterized PDEs, retrieving 215.7 million putative PDEs [cite PETadex v1]. We demonstrated that novel PDEs can be discovered in this dataset through operationally random selection. However, the lack of annotations prevents strategic selection of candidates for validation.

Here we present PETadex, a methodological and quantitative rebuild of our previous dataset that supports the targeted selection of PDEs. Using 411 representative PDEs as input, we searched NCBI-nr and the SRA (via the Logan assembly), filtering hits with catalytic-residue-constrained profile HMMs. The resulting dataset of [x] million enzymes carries catalytic domain boundaries and is annotated with fold assignments, predicted protein structures and the temperature and pH of the biosamples of origin. Candidates can therefore be filtered by catalytic integrity, structural fold and environmental meta-data before testing, enabling high-throughput targeted selection that can be further automated with machine learning. The enzymes, annotations and a toolkit of methods are freely available at <https://petadex.net>.

Results

Sequence and structure-based annotation of 609 known plastic-degrading enzymes

Following established workflows for the homology-based retrieval and expansion of a given set of enzymes of interest [cite Erickson, Danso, Logan etc.], we built a workflow for the retrieval of PDEs from NCBI-nr and the SRA (using the Logan assembly).

Improving upon our previous PETadex dataset, by updating the input search query and fixing several methodological limitations, we present the updated PETadex v2 dataset.

To generate this updated PETadex-260701 dataset, we created a starting search query from combining two versions of the Pazy database, the first from the original PETadex generation (pazy-240826; $n = 212$; August 26th, 2024), and a recently updated version (pazy-260701; $n = 473$; July 1st, 2026). Sourcing 136 enzymes from pazy-240826, 397 enzymes from pazy-260701 and 76 merged between both, generating a 609 enzyme union set (plasticome). Plasticome was then clustered at 90% amino acid identity to 411 cluster centroids (Fig. 2A). Predicting structures of these centroids using ESMFold2 (ESMC-6B trunk),

Fig 2A. Plasticome Composition



Fig 2B. PETadex Generation




Fig 2C. PETadex Clustering




Fig 2D. Cluster Size Distribution




Fig. 2 The PETadex-260701 dataset. (A) Composition of the plasticome: the 609-enzyme union of the two PAZy snapshots, its collapse to 411 cluster centroids at 90% amino acid identity, the 46 graph components those centroids resolve into, and the CATH folds assigned to them, coloured by provenance. (B) Overview of PETadex generation, from the 411 input sequences through the PETadex-nr and PETadex-logan expansions. (C) Hierarchical clustering of catalytic sequences into enzymes (90% identity), families (60%), superfamilies (30%) and distinct folds. (D) Cluster size distributions at the protein, enzyme and family levels.

then searches against the CATH S40 database with Foldseek alongside searching against InterProScan resolved the proteins into 15 distinct protein domains. These centroids are further aligned into 46 graph components ($\geq 30\%$ amino acid identity and $E < 1 \times 10^{-5}$) (Fig. 2A).

Homology-based retrieval yields a dataset of 215.7 million putative plastic-degrading enzymes

To expand the dataset of known PDEs beyond the plasticome set, we first searched these 411 centroids into the NCBI nr database using DIAMOND(DIAMOND v2.2.4.184, `--very-sensitive, -k0`), recovering 42,053,081 pairwise alignments, collapsing to 8,850,583 unique sequences [cite]. Of these, 7,825,535 pass the e-value filter ($E \leq 1 \times 10^{-5}$), and further

removing plasticome self-hits returns a final count of 7,825,107 sequences (Fig. 2B).

In order to search Logan with only the catalytic domains of these plasticome homologs, we built profile HMMs to characterize the plastic-degrading domain boundaries and catalytic residues. These HMMs were built from seed MSAs (MUSCLE 5.3) of plasticome sequences together with their closest nr homologs ($E \leq 1 \times 10^{-50}$), where each member of the alignments must contain the residues involved in the catalytic mechanism, annotated from literature [cite]. The 46 graph components were further segmented using these catalytic mechanisms, resolving into 53 MSAs. The HMMs were built from the seed MSAs using a pipeline analogous to JACKhmmer [cite]. HMMs were created using HMMbuild, and then searched (HMMsearch) against close NCBI-nr homologs ($E \leq 1 \times 10^{-10}$), annotating the catalytic residues [cite]. Alignments were created from sequences that were below the e-value cut-off of 1.0E-10 and contained all of the canonical catalytic residues. These MSAs serve as the seed MSAs for the next iteration of the pipeline. While creating the next iteration’s HMMs, match and insert state annotations were propagated from the previous HMM and any column of the alignment that had an occupancy below 0.7 was declared an insert state. The pipeline was then repeated for a total of 4 iterations.

Searching these component HMMs against the 7,825,107 nr sequences yielded 7,562,397 HMM hits. Of these hits, the HMM range of 6,857,582 contained all of the positions of the catalytic residues (“complete hits”), of which 6,343,582 contained the canonical catalytic residue identities. The ranges annotated as the catalytic domains were then clustered (DIAMOND –linclust), resulting in 2,740,458 90% centroids (Fig. 2B). These centroids serve as the input query into Logan.

Expanding on the nr search, we queried these [x] million enzyme homologs across Logan (version [x]) contigs, recovering [x] billion matching sequences (Fig. 2B). Of these Logan sequences, [x] matched an nr enzyme at 90% identity, and clustering the remaining [x] at 90% aaid, yielding [x] million non-redundant enzyme homologs (Fig. 2C,D). We’ve dubbed this dataset PETadex-v2, with sub sets dubbed PETadex-plasticome (or

simply plasticome), PETadex-nr and PETadex-logan, depending on the origin. [Describe a domain from the search, can be a/B hydrolase, Amydases, etc.]

Structural and taxonomic annotation of the PETadex enzyme space

To explore the PETadex dataset, we annotated its sequences structurally and taxonomically, and linked each sequence to the metadata of the sample it was recovered from.

We first predicted structures for the 22,235 superfamily centroids (30% amino acid identity) using ESMFold2 (ESMC-6B trunk). These structures resolve both the catalytic domain and the accessory domains flanking it, which may additionally contribute to the catalytic mechanism.

We next assigned each sequence a source organism, from the NCBI taxonomy of its GenBank record (PETadex-nr) or the taxonomic composition of its source run (PETadex-logan). Of [x] PETadex-nr sequences, 1,192,350 ([x]%) carry an assignment, overwhelmingly bacterial (1,071,209; 89.8%), followed by Eukaryota (103,380; 8.7%), Archaea (17,493; 1.5%) and Viruses (268; <0.1%) [taxonomy panel removed from Fig. 3; cite a Supplementary figure or drop the reference].

Sample metadata provides planetary-scale annotation of the PETadex dataset

Each PETadex-logan sequence is taken from an SRA run, which links to an associated BioSample record, and potentially to an overarching BioProject. These records carry metadata features, such as the location, biome, host, etc., describing the sample, its collection conditions and subsequent sequencing methods. For our dataset of 307,155,746 ORFs, 256,236,322 are linked to BioSamples (83.4% coverage), drawn from 5,170,787 BioSamples (Fig. 3A).

Tracking the coverage of the metadata fields for the BioSamples, we can see that the BioSample records are most complete when describing collection date (47.5%), stated coordinates (32.0%) and isolation source (28.4%) (Fig. 3A). From the geographical coordinates of the samples, we can visualize the truly global span of PETadex




Fig. 3 Environmental metadata of the PETadex corpus. (A) Share of corpus ORFs ($n = 307,155,746$) with each metadata field. Solid bars, stated in BioSample records; hatched bars (*), inferred from BacDive optima of the taxa in each BioSample. (B) PETadex sequences by stated BioSample coordinates (log scale). (C) Environment categories of profiled BioSamples (left; $n = 2,478,477$) and of 90% clusters with ≥ 2 profiled ORFs (right; $n = 6,904,458$). Blue, environmental; grey, host-associated or other.

putative PDEs, with notable hotspots in North America, Europe and East Asia. It is important to note, however, that the span of these enzymes largely tracks sequencing runs globally, or simply put, these enzymes are found anywhere and everywhere sampling occurs. Other metadata features describing the physical conditions at collection, such as temperature and pH, are rarely documented, with stated temperature available for 2.1% of ORFs and stated pH for 0.7%. This reflects the sparse and often unreliable nature of physical condition measurements throughout the SRA.

BioSample records are also highly heterogeneous, for instance the isolation source field alone contains 78,439 distinct values. To make these

records comparable, and to allow enzyme candidates to be selected for experimental validation by their environment of origin, we classified BioSamples into 21 environment categories based on their stated environment terms, isolation source and host (Methods). Of the 2,744,599 BioSamples carrying at least one usable metadata field, 2,478,477 were assigned a category. Through their BioSample, 52.4% of the 307 million ORFs receive a category directly. To extend this coverage, we assigned each 90% sequence cluster the dominant category of its member ORFs and propagated it to members without a category of their own, raising coverage to 97.1% of ORFs (Fig. 3A). Counted by BioSample and by cluster, the categories give opposite pictures (Fig. 3C). Human-associated

categories account for 54.3% of profiled BioSamples, predominantly *human other* (31.1%) and *human gut* (11.6%), but only 1.7% of clusters with at least two profiled ORFs. Environmental categories account for 18.7% of BioSamples but 82.2% of clusters, led by soil (31.1% of clusters) and followed by other environmental sources (16.6%), marine and freshwater (8.4% each), sediment (7.2%) and wastewater sludge (6.7%). Thus, while sequencing effort in the SRA is dominated by host-associated samples, most of the sequence diversity of PETadex enzymes lies in environmental samples.

To improve the sparse metadata annotation for environmental physical conditions of temperature (2.1%) and pH (0.7%), we attempt to leverage community species composition for BioSamples. For each run, we first averaged BacDive growth optima of the taxa detected and then weighted by read abundance to obtain an approximation of sample conditions (Methods) [cite BacDive]. This proved effective at raising coverage from 2.1% to 80.6% of ORFs for temperature, and from 0.7% to 65.8% for pH (Fig. 3A). These inferred optima describes the organism within a sample, so we validated them against stated measurements for BioSamples where data was available. Abundance-weighted optima correlate with measured pH across 6,473 BioSamples (Pearson $r = 0.34$, against 0.14 without weighting) and with measured temperature across 45,741 BioSamples ($r = 0.16$, against 0.05). These samples are compressed towards neutral pH and mesophilic temperatures, and therefore do not improve on a constant in absolute error. We therefore propose that inferred optima can be applied as a means of ranking and filtering samples, but is not reliable as an estimate of *in situ* conditions, and should thus be treated as distinct from stated values.

Together these annotations support homology-based characterization of uncharacterized enzymes, and allow candidates to be filtered on catalytic integrity, fold and habitat prior to testing.

Website

This homology-based characterization pipeline is implemented through a DIAMOND sequence search interface available freely at petadex.net. This sequence search runs 20 shards in parallel

to achieve an average search time of ~ 47 seconds over the database of ~ 215.7 million enzymes.

Discussion

Methods

Supplementary information.

Acknowledgements.

Declarations

- Funding
- Competing interests
- Ethics approval and consent to participate
- Consent for publication
- Data availability
- Materials availability
- Code availability
- Author contribution

References